

Indexes

Agenda

- ① what are indexes
- ② How indexes work
 - ↳ B & B⁺ tree
- ③ Cons of indexes
- ④ Indexes on multiple columns
- ⑤ Indexes on strings
- ⑥ Indexes PRACTICAL
 - ↳ EXPLAIN.

e.g. for join

for i — 100; $\frac{10^9}{10^9}$
for j — 100;
doSomething(); ; } joins

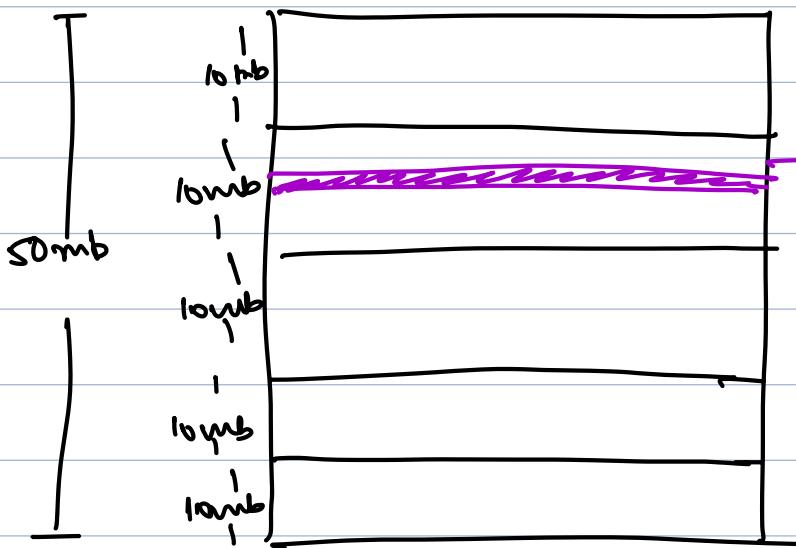
$$10^9 \times 10^9 = \underline{\underline{10^{18}}}$$

if the dB exactly was doing this

↳ Performance will be very slow.

- ① dB have to do a lot of optimizations to make queries fast.
- ② dB stores data on disk

Hard Disk



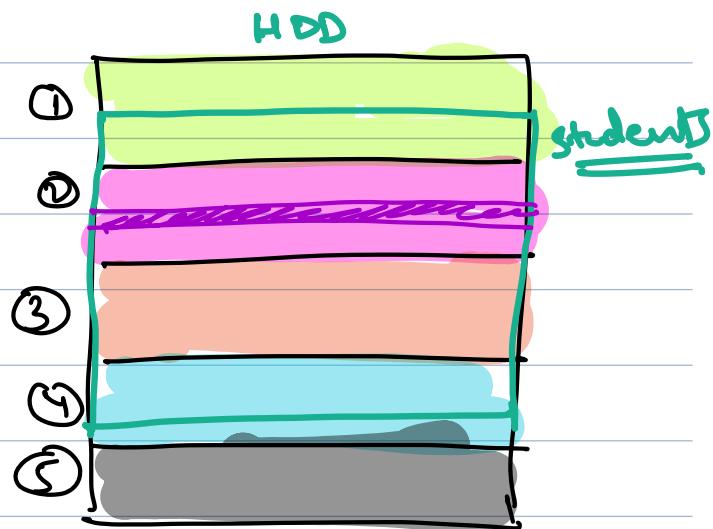
each block has data.

our data is saved on some block.

OS gives the memory complete block that contains the data.

select *
from Students
where id = 100;

- ① Bring block 1 to memory
- ② Check all rows if they contain $id = 100$



Not OPTIMAL.

⇒ unnecessary blocks that are being fetched are a waste of performance.

Table of contents.

SNO	title	P.NO
-----	-------	------

- Index in a book helps find page faster
- Index in a DB → helps find blocks of disk faster



block / blocks that
contain the rows
that we define.

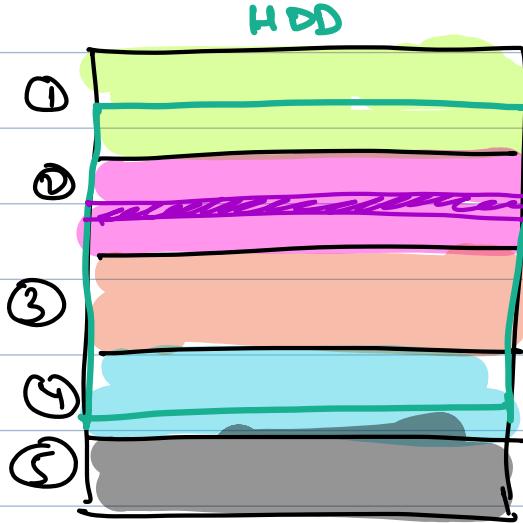
Purpose of indexes: Reduce no. of disk block
accesses to fetch data.

How do indexes work..

Huge table 100M records/rows
Select * from students where id = 100;

↓ index ↗

<u>id</u>	<u>block#</u>
1	1
2	1
3	2
4	2
5	3



Map | Hashmap

(key / value)
(fast)

without hashmap → 4 blocks

with hashmap → 1 block

Q Do indexes work on non-unique columns?

select * from students
where name = 'Ujjwal';

students

<u>id</u>	<u>name</u>	<u>bsp</u>
1	Ujjwal	100
2	Meera	80
3	Ujjwal	78

Index

<u>name</u>	<u># block</u>
Ujjwal	[1, 3]
Meera	[2]
Ajaay	[4, 5, 6]

HM < String, List<block no>>

select * from students

where name = 'Ujjwal';

—

1 ✓
2 ✗
3 ✓
4 ✗
5 ✗
6 ✗

without index → 6

with index → 2

where $x = y$;

⇒ there are also queries about ranges.

e.g. find all the students with name $\geq 'M'$
and name $\leq 'O'$;



start with M, N, O —

find all students with fsp blw 60 & 70.

index

fsp	block
42	5
66	2
100	8
22	1

keymap < int, string)

get all keys blw
20 & 40

for i = 20 → 40 :

if i in map :
map.get(i);

with a hash map it is very tough to get all values within a range.

↓

not optimal

Hashmap - check a value $\rightarrow O(1)$

get all values in a range $\rightarrow O(N)$

range query like $[l \dots r]$

① Sorted

② Key : Value

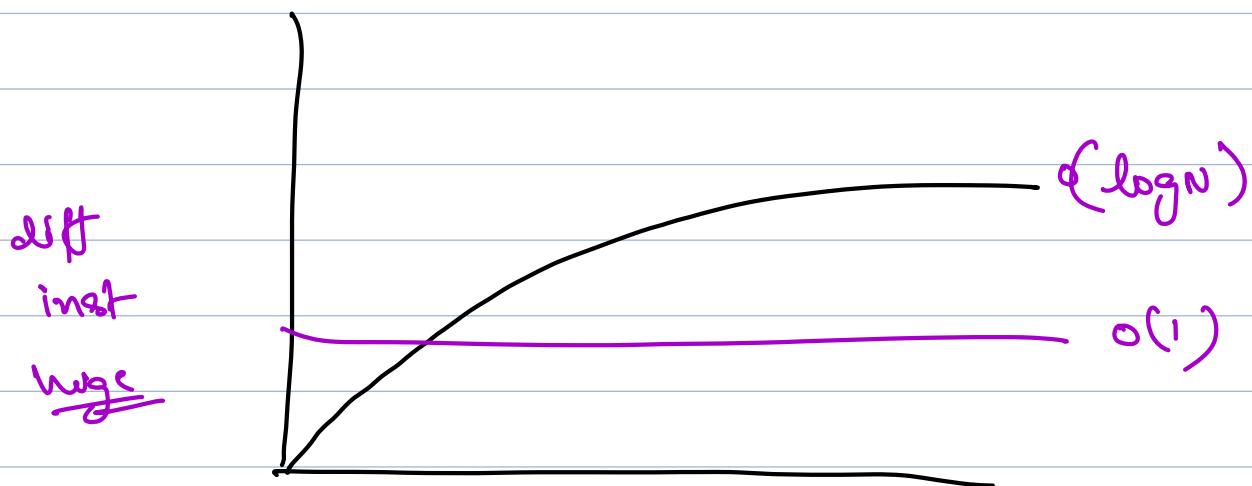
Sorted

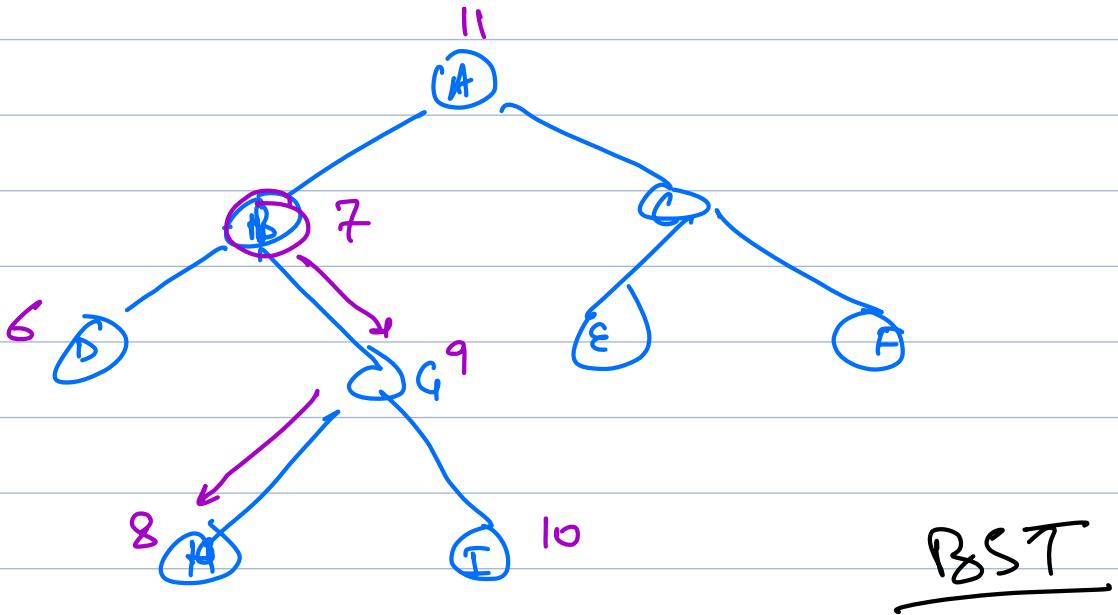
sorted
map

→ Treemap.

writes on
Balanced BST

check a value $\rightarrow O(\log N)$





Given a node
get the next bigger node?

$$B \rightarrow J \quad O(\log N)$$

$$40 \rightarrow 60$$

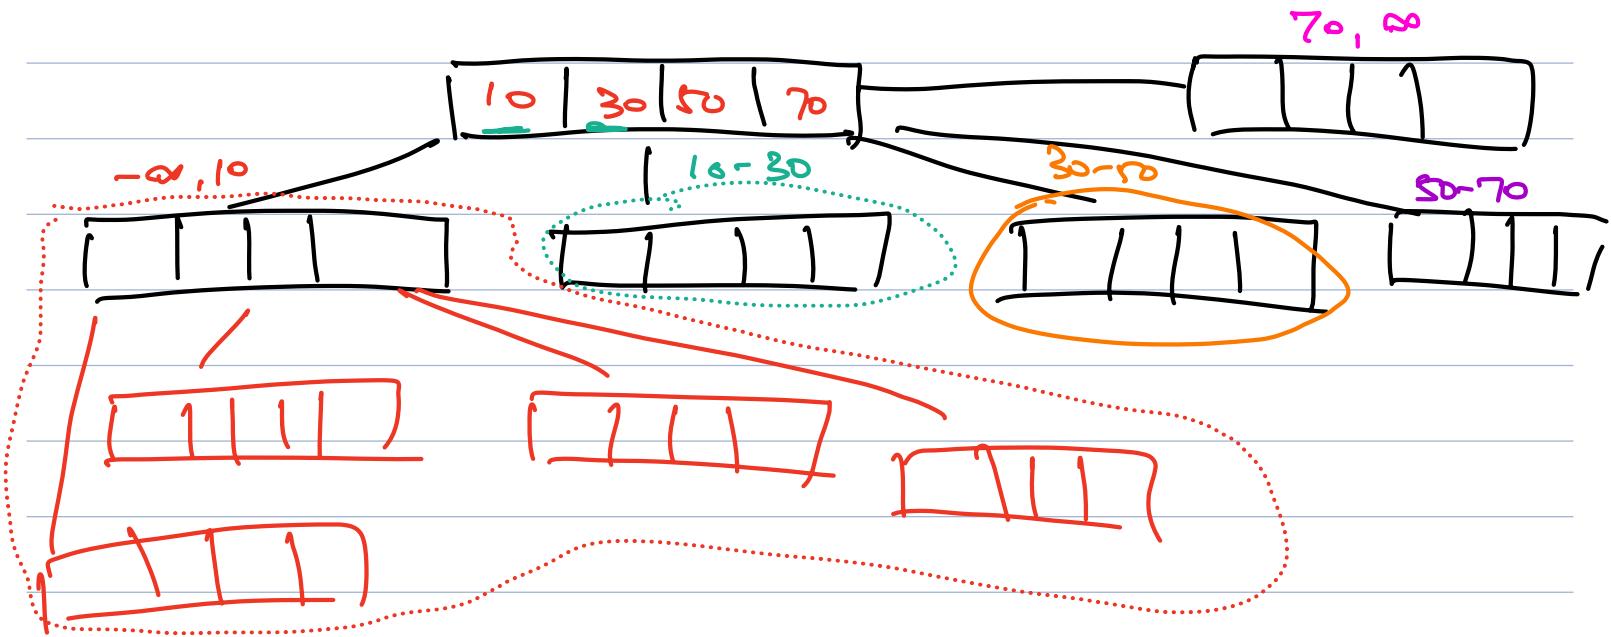
- ① go to 40
- ② keep going to next node till
you reach 60

Indexes work very similar to TreeMap

This method is used to reduce no. of disk access to increase performance of SQL query.

Indicies use data structures called B/B⁺ trees;

→ instead of every node having ≤ 2 children.
every node can have $\leq m$ children.



Because a B tree has many children

→ The height of B tree will be lower.

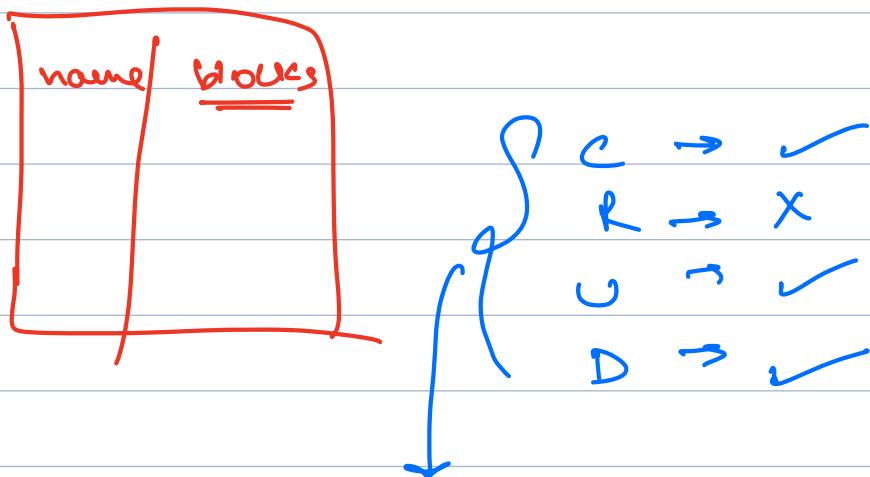
$O(H)$

Checking a Value $\rightarrow O(H)$ & faster
getting a range $\rightarrow O(H)$

Are you clear with:

- ① How indexes work
- ② What is benefit of index
- ③ Why hashing is not good

CONS of indexes.



• Write operation
(C, U, D) → changing data on disk

① Writes are slower.

→ Typically a copy of index is also stored on disk.

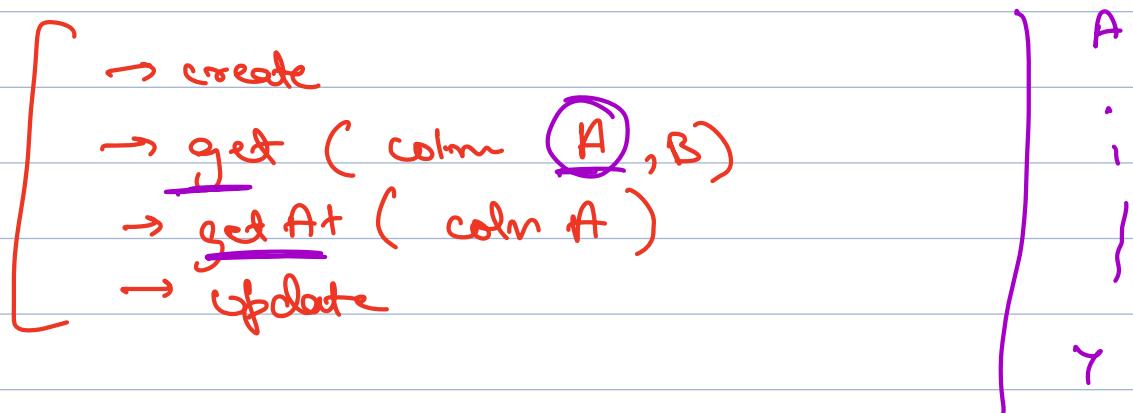
② Storage requirement increases.

* Don't create index prematurely.

→ Create an index only when you see the need of it.

[Break]

08: 30 AM.



Indexes on multiple columns.

(id)
(name)

(name, psb)
(name, psb, rank)

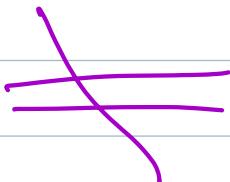
→ We can have index on multiple columns.

Index \rightarrow (name, b_{sp})

query select *
 from students
 where b_{sp} = 80;

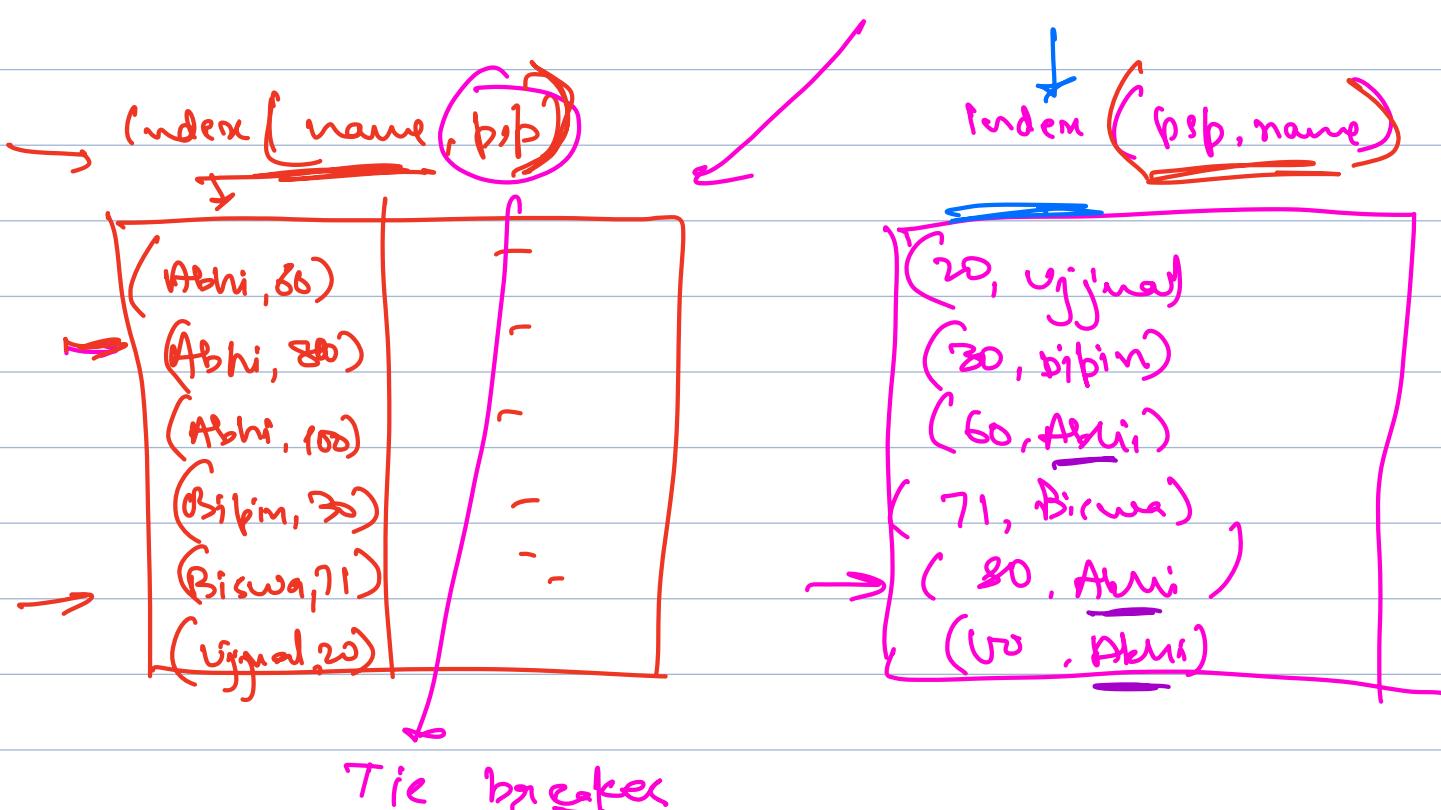
Will this index help? \rightarrow NO

If we create an index on (name, b_{sp})



Creating index on (name)
 +

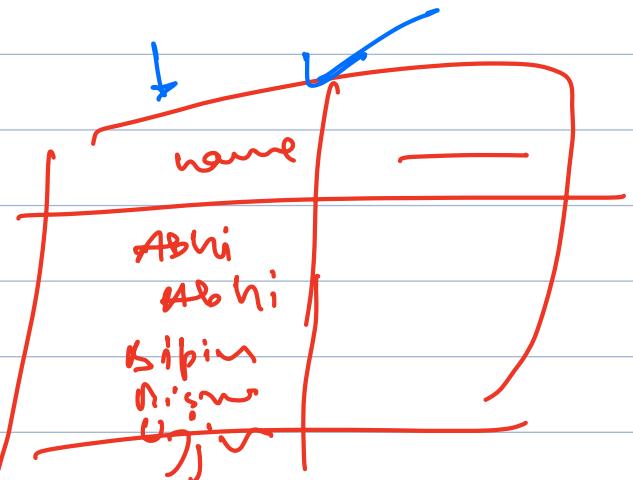
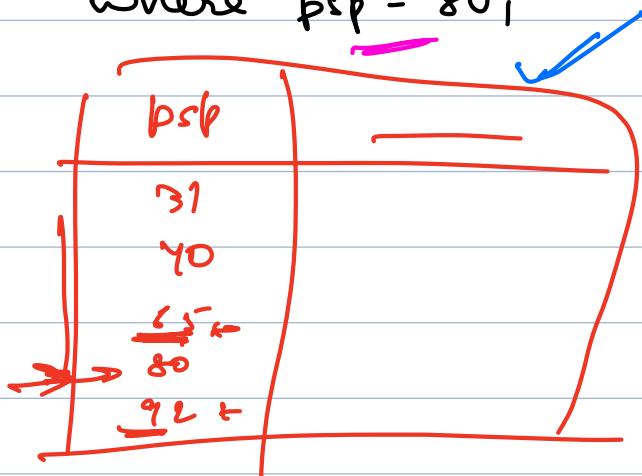
Creating index on (b_{sp})



select *

from students

where bsp = 80;



Scenario:

query	index on	Help	Not Help
① name	bsp		✓
② name	(name) (bsp)	✓	
③ name	(<u>bsp</u> , name)		✓
④ name	(name, bsp)	✓	
⑤ name = x and bsp = y	(bsp, name)	✓	

* If a query is on coln x

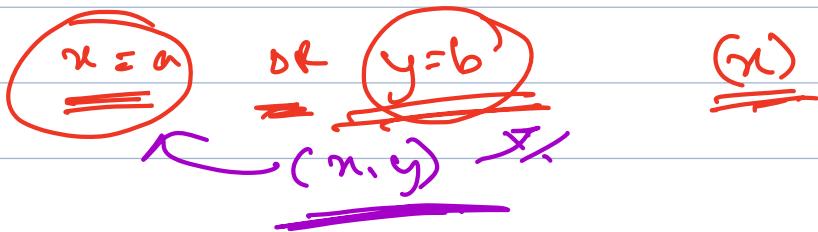
any index with x in the prefix will help.

- (x, a, b, c) ✓
- (x, a) ✓
- (a, b) ✓
- (a) ✓
- (a, x) ✗

Scenario

$x=a$ and $y=b$

Index
 $(x) (y)$ —
 $\underline{(x, y)}$ —



AND query multiple column index will help
 OR query multiple column index WON'T help.

Indexes on Strings

Users		
id	name	email
		=

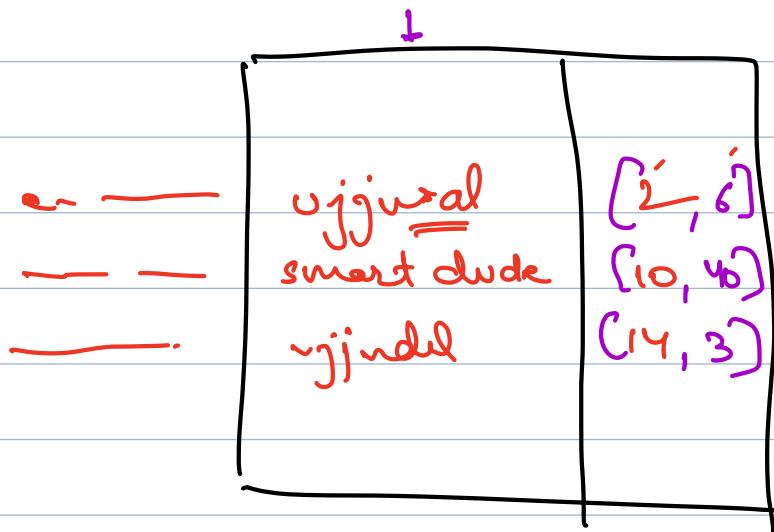
select *
from users
where email = ' ';

email	block #
ujjwal@scaler.com	2
smartduke@gmail.com	10
ujjinder@gmail..	14

- ① size of the index will be huge.
- ② String matching will be slow.

How do we solve this?

Instead of creating an index on complete email,
we create index only on part before '@'.



→ Size of the index is reduced.

① Space is saved.

select *
from users

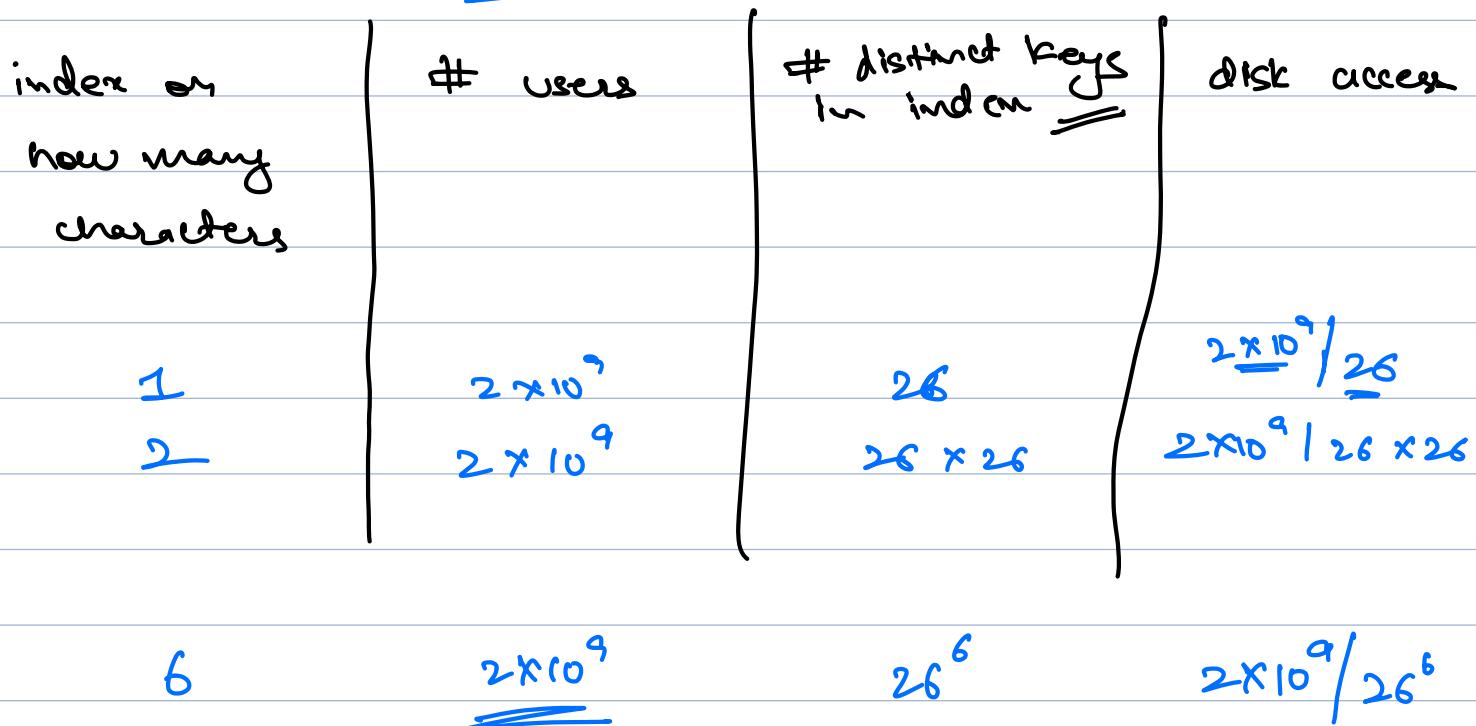
where email = 'ujjwal@scaler.com'; \Rightarrow will have to
fetch 4 blocks
instead of 1 block
earlier.

[without index \rightarrow all block]
with complete index \rightarrow 1
with partial index \rightarrow 4

Typically when you have string columns,
index is created on a prefix of the column
instead of entire column.
↓
enough mostly.

Address → Index address[:r];

$$\text{Google} \rightarrow 2 \times 10^9$$



Q select *

from users

where addresses like '%. toronto%';

An index will not help here.

Full Text Index.