

Transactions - I

Agenda:

- ① Transactions
- ② Properties of a transaction

→ A

→ C

→ I

→ D

Imp for
interviews

- ③ Read and write op[“]
- ④ Commit and Rollback
- ⑤ Transaction Isolation level

→ lost updates

→ dirty reads

→ Phantom reads

→ Non-repeatable reads

Imp for
interviews

- ⑥ Deadlock

A bit more complex & theoretical class.

What are TRANSACTIONS

Till now we have used SQL for a single query. What if we need to perform a group of operations to achieve a task?

Bank Example

Transfer money

₹ 1000



(A)

₹ 5,000



(B)

for this transaction / transfer to happen what changes do we need in the accounts table.

Accounts

id	name	balance	created-at	branch-id
1	A	₹ 1000	-	-
2	B	₹ 5000	-	-

- ① Check balance of A → DB call → Read
- ② Check > 500 ⇒ Application code
- ③ Reduce balance of A by 500 → DB call - Update
- ④ Increase balance of B by 500 → DB call - Update

transferMoney (from, to , amount) {

 if(bal >= amount) {

 } else {

}

do you think only one person will be calling this function at a point of time?

① transferMoney may be getting executed by multiple people at the same time.

A → B (₹500)

C → B (₹500)

eg ecommerce - Amazon

- | | | |
|---|-----------------------|----------------|
| ① | Check balance of A | Read A |
| ② | Check > 500 > | Not read/write |
| ③ | Reduce balance of A | write |
| ④ | Increase balance of B | read & write |

B + = 500 \Rightarrow This is NOT a single operation.
↓

B = B + 500 \rightarrow first some has to read the previous value of B.

↓

} Read B → temp
 temp → temp + 500
 Write B ← temp

① A B 500

② C B 10000

transferMoney (a , b , amount) ?

If database
went down
after this point?

A → 1000 500

B → 5000 5500 15000

C → 15000 5000

Read A ⇒ x ① → 1000 ② → 15000

If ($x >= \text{amount}$) {

 write A ← x - amount ① → 500 ② → 10000

 Read B ⇒ x ① → 5000 ② → 5000

$x := x - \text{amount}$

 write B ← x ① → 5500 ② → 15000

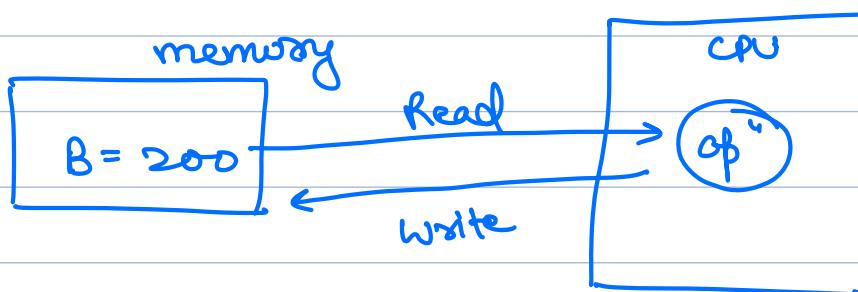
21000
Initial sum

20500
final sum

}

Two operations
can't happen at
the exact same
time. It will be
one after the
other.

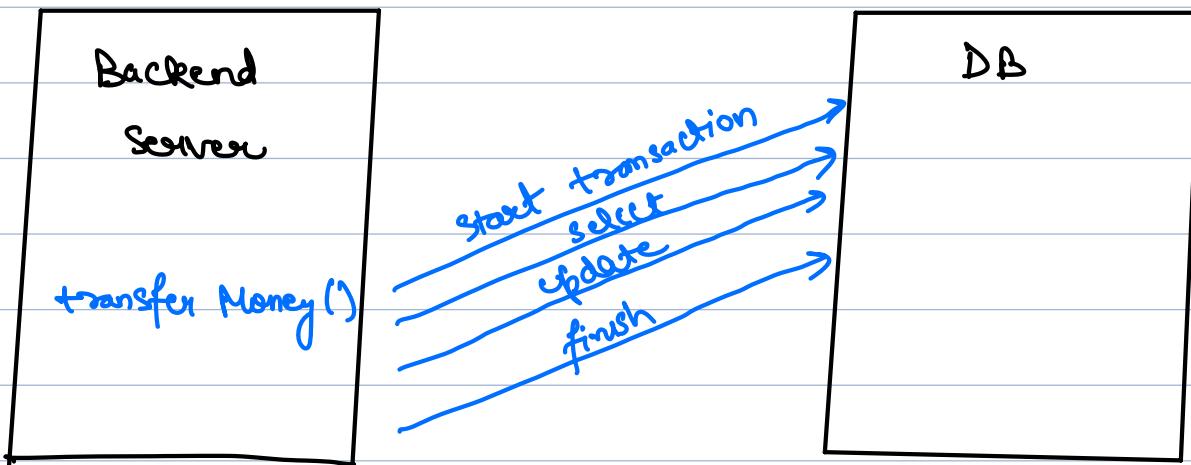
How a system works



- ① Inconsistent / illogical state
- ② Entire operation may not complete

Transaction — A set of DB operations logically grouped together to perform a task.

Transfer money:



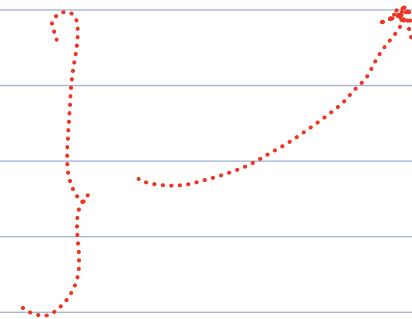
As we learnt when we group a lot of operations, a number of things can go wrong.

Expectations | Properties of a Transaction.

#Properties that we MAY want from our transactions at times.

(ACID Properties)

- A → Atomicity
- C → Consistency
- I → Isolation
- D → Durability



ATOMICITY

↓
from the word Atom → Smallest **in divisible** unit of mass.

⇒ Transaction should appear indivisible to user.

→ to user outside it should feel that either nothing has happened or everything has happened.

[behave like ONE single unit / operation]

→ A transaction should never end in an intermediary state.

→ never the DB in between



Kya be?

CONSISTENCY

→ correctness

→ exactness

In the above example atomicity was handled → Transaction either happened & money transferred or nothing.

But work was not correct. Final outcome was wrong.

→ Accuracy

→ Logical correctness

Example where consistency might not be needed.

	Count → 1.6 cr	
①	16000000	↖ B
②	16000001	16000000

Hotstage	
Live IPL match	
stream_id	Count
—	—

- ① get current count
- ② update count

consistency hampered
BUT performance improved

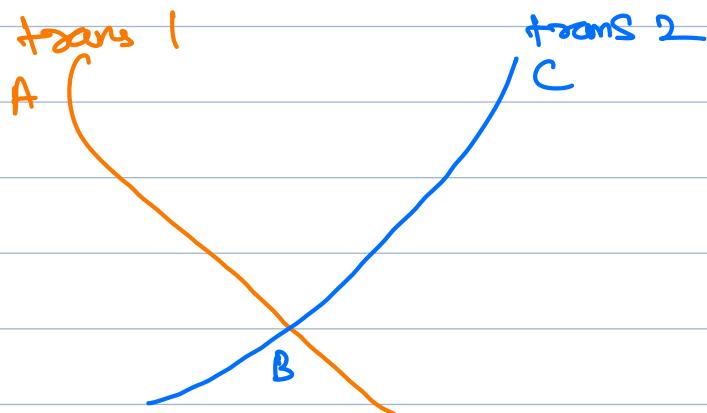
HAPPY INDIA .

The state of DB before "trans" started and also after "tran" completed should be accurate.

ISOLATION

→ One transaction shouldn't affect other transaction running at the same time running on the same data.

In the above case, we wouldn't have any issue if only one transaction was taking place. The 2 transactions were not isolated.



If you have 2 gf/bf if they don't interfere, you will not have any issue.

DURABILITY

Once a transaction is completed we would want its work to stay persistent.
→ store it on disk.

Seems obvious but may not be needed.

You might think when you do write operation to DB, it instantly writes to disk. But NO

In case where DB has to provide fast writes, it might do something like a batch write.

↓
durability lost

→ Ques can be what if disk gets corrupted

→ Data centre catches fire

→ Alien invasion .

} Different levels of durability

[BREAK]

Commit and Rollback

Students

id	name	psp	batch-id
1	Ujjwal	80	2

update student

set psp = 100

where id = 1

select * from students .

↓

we expect psp=100.

This is based on assumption , whatever we did in the first trans" worked.

This worked because our query's result got stored/saved to the dB.

WHY

Autocommit:

- every SQL query starts a transaction.
- executes itself.
- saves the changes // Commit

↓

What is commitment?

- # → Ask commitment from boss for bonus
→ Committed relationship

① Start transaction

② Execute operation

③ Commit

↪ Statement that is used to persist the results of an SQL query.

SQL query only makes changes but won't commit them to the database.

These are saved only after commit.

* Commit takes care of durability.

Autocommit = true

By default

[code] (make two sessions)

- set autocommit = 0;
- select * from film where film_id = 10;
- update film
set title = 'Vijawal'
where film_id = 10;
- commit;

Workbench takes snapshots of database per session to achieve this.

→ → →

If a lot of people are transferring to B, A → B, C → B etc and we suspect B's account, we would want to block all transactions.

So we UNDO money reduced from A's account.

ROLLBACK

rollback;



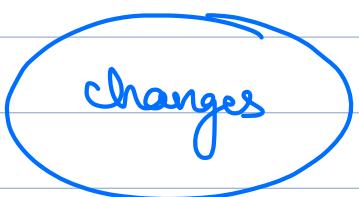
undo all changes done since the last commit.
(changes still in the memory)

commit → marriage → persisting the changes to dB.

rollback → break up → undo whatever changes were being done since the last commit.

→ You cannot rollback after a commit is done.

Session 1



Session 2



these changes weren't visible to Trans' 2.

Both sessions are seeing diff things

Why? MySQL → each session is isolated

4 levels of isolation

① Read uncommitted

PostgreSQL → ② Read committed

→ ③ Repeatable read

Default isolation level ④ Serializable

for MySQL

most relaxed

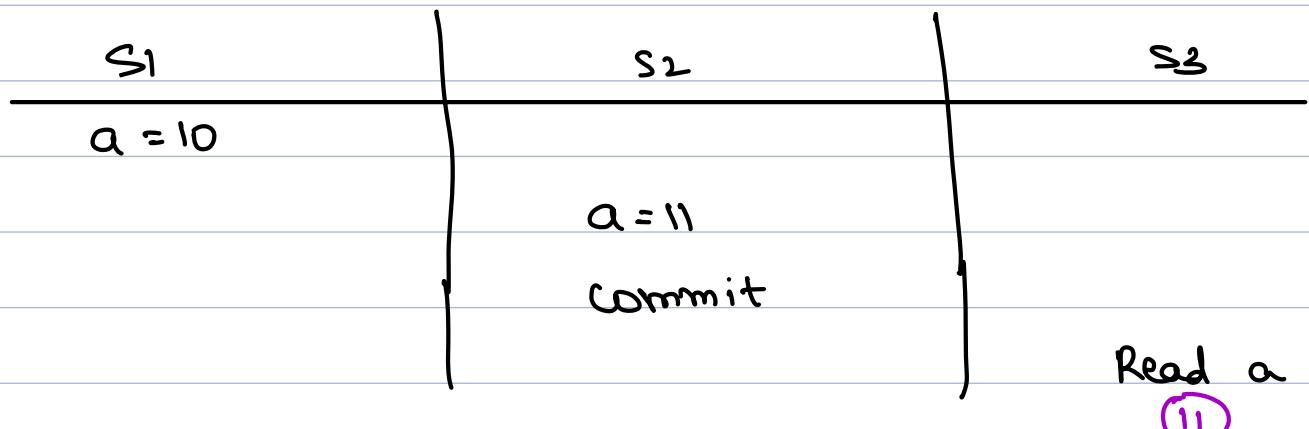
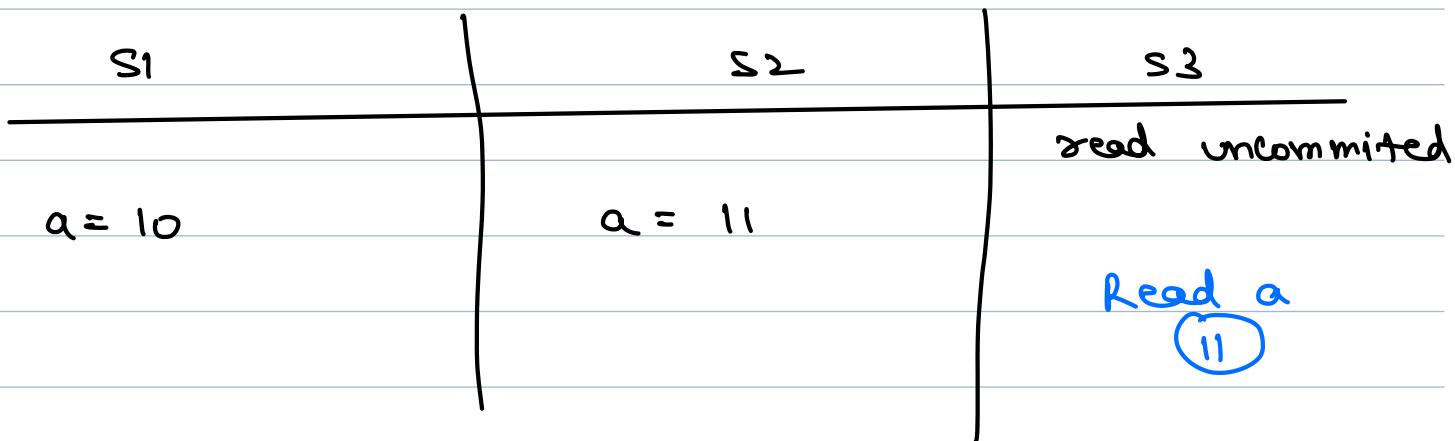
Increasing
order
of
severity

most strict

* Isolation levels we're talking about what data is read, not what is updated.

Read uncommitted

- Allows a transaction to read even uncommitted data from another trans.
- latest uncommitted data (committed or uncommitted)



★ Isolation level only talks about how your session will act (read data).



The isolation level of other trans doesn't matter.

[code]

- Show variables like 'transaction-isolation';
- Set session transaction isolation level read uncommitted;

PROS.

- ① Fast → performance

CONS.

Trans1 (RR)

$A = 2000$ (2090)
 $B = 2000$ (1900) Money lost

- ① Read $A \rightarrow x$ (2000)
- ② $x \leftarrow x - 10$ (1990)
- ③ Write $A \leftarrow x$ (1990)

- ⑦ Read $B \rightarrow x$ (^{2000 if RE}
~~1900 if RU~~)

$x \leftarrow x + 10$

Write $B \leftarrow x$

Commit ;

Trans2. (RU)

- ④ Read $B \rightarrow x$ (2000)
- ⑤ $x \leftarrow x - 100$ (1900)
- ⑥ Write $B \leftarrow x$ (1900)

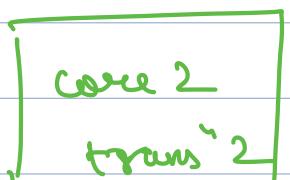
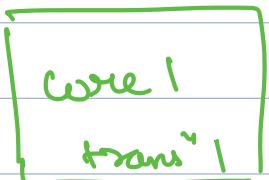
- ⑧ Read $A \rightarrow x$ (1990)

- ⑨ $x \leftarrow x + 100$ (2090)

Write $A \leftarrow x$ (2090)

Commit ; ($A \rightarrow 2090$)
 $(B \rightarrow 1900)$

Transactions happen on CPU, who owns lines
 one by one with context switches [Multi core machines]



Because 2 diff "trans" were running in parallel



ended up in inconsistent state

RU hampered consistency.

At line 8 \Rightarrow "trans" 2 read data that was not committed and rolled back.

DIRTY READ

(Risky read)

when a "trans" reads data that is not committed.

That data might change | rollback | not get committed