

## Transactions - II

### Agenda

- 1) Read uncommitted ✓
- 2) Read committed
- 3) Repeatable read
- 4) Serializable

Isolation levels.  
(Locks)

### Deadlocks

Dirty read - Reading some data that is not yet confirmed yet.

# How to solve this?

A data is confirmed once it is committed.

Read committed

→ Your trans will only read the latest committed data.

[code]

L set session transaction isolation level read uncommitted.

R set sessi transaction isolation level read committed.

L `set auto-commit = 0;` R

`select * from film where film_id = 10;`

update film,

`set title = 'ujjwal 4'`

`where film_id = 10;`

# datagrip

# Transaction isolation level of yours matters.

DON'T care about others isolation level.

R `select * from film where film_id = 10;`  $\Rightarrow$  get latest committed value

# When you start a transaction, if turns off the auto commit till the transaction is completed.

R update film,

`set title = 'ujjwal 6'`

`where film_id = 10;`

$\Rightarrow$  It will go to waiting bcoz there is a lock taken on the other side.

L `commit;`

Read  $\Rightarrow$  Shared lock

Written  $\Rightarrow$  Exclusive lock

Some one has  
Shared lock

someone else can  
⇒ read  
⇒ write

Exclusive lock

↳ given at a  
write query

→ read  
→ write X

# Read on left side in code & see right side  
can write as left one takes a shared lock.

only

new that was inserted |  
deleted/updated.

If one transaction  
transaction will be allowed to write on that  
row.

# focus on isolation levels,  
we talk about locks  
later.

Read committed

- A trans' will read always the latest committed value
- never read an uncommitted value.
- No dirty reads.

## Problems with this.

Users			
id	name	psp	email_sent
1	A	60	false
2	B	<del>80</del> 79	false
3	C	70	false
4	AB	45	false

Q Send email to all students with psp < 80.

list <users> = select \* from users where psp < 80;  
[1, 3, 4]

send emails

⇒ At this point  
2's psp goes to 79.

what  
can  
go  
wrong?

update users

Set email\_sent = true

where psp < 80;

[1, 2, 3, 4]

# This is wrong as user 2 did not get an email.

Trans 1

① Read users with psp < 80

② doing work

Trans 2

③ update psp of user 2

④ Read users with psp < 80

# Within a trans , the value / result of a particular query can change.

⇒ Within a transaction if I read the same row again , it must have the same value.  
↓

# When a trans starts, it should take a snapshot of the db and only that value should be taken.

↓

Read Committed will not have repeatable reads.

+

It has non-repeatable reads.

# In read committed ; it always reads the latest committed value , even if that committed value happened after the trans started .

[CODE]

RC

Update

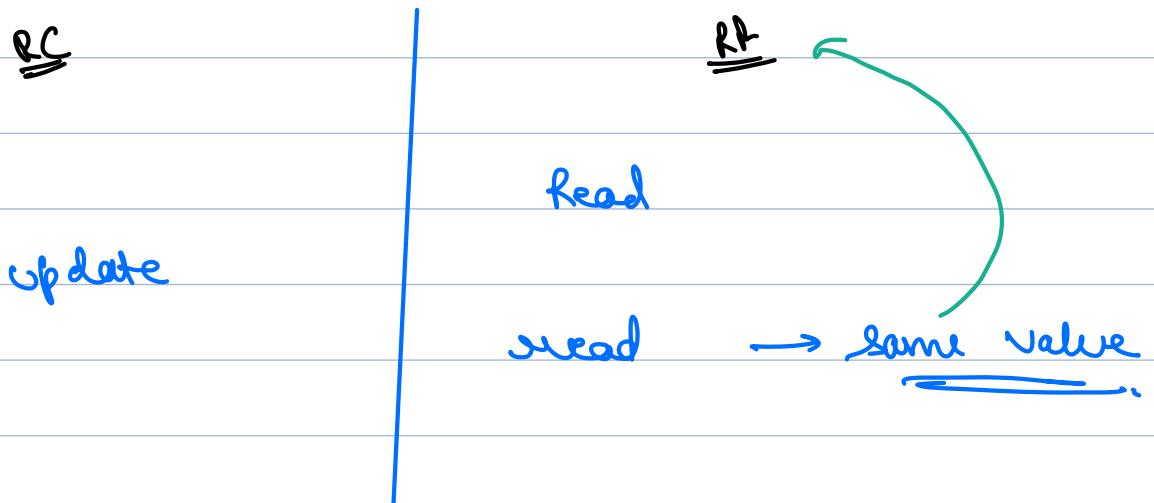
RC

Read

Read → diff value from above .

# make sure autocommit is off.

Set session transaction isolation level repeatable read;



# This will not work if autocommit is ON as after every statement it will commit & start a new transaction.



Repeatable reads. works only in the scope of a single transaction.

autocommit = 0 :



① write a query  $\Rightarrow$  starts a new trans

② commit  $\Rightarrow$  end the trans.

autocommit = 1

Start trans  $\rightarrow$  executing  $\rightarrow$  end trans

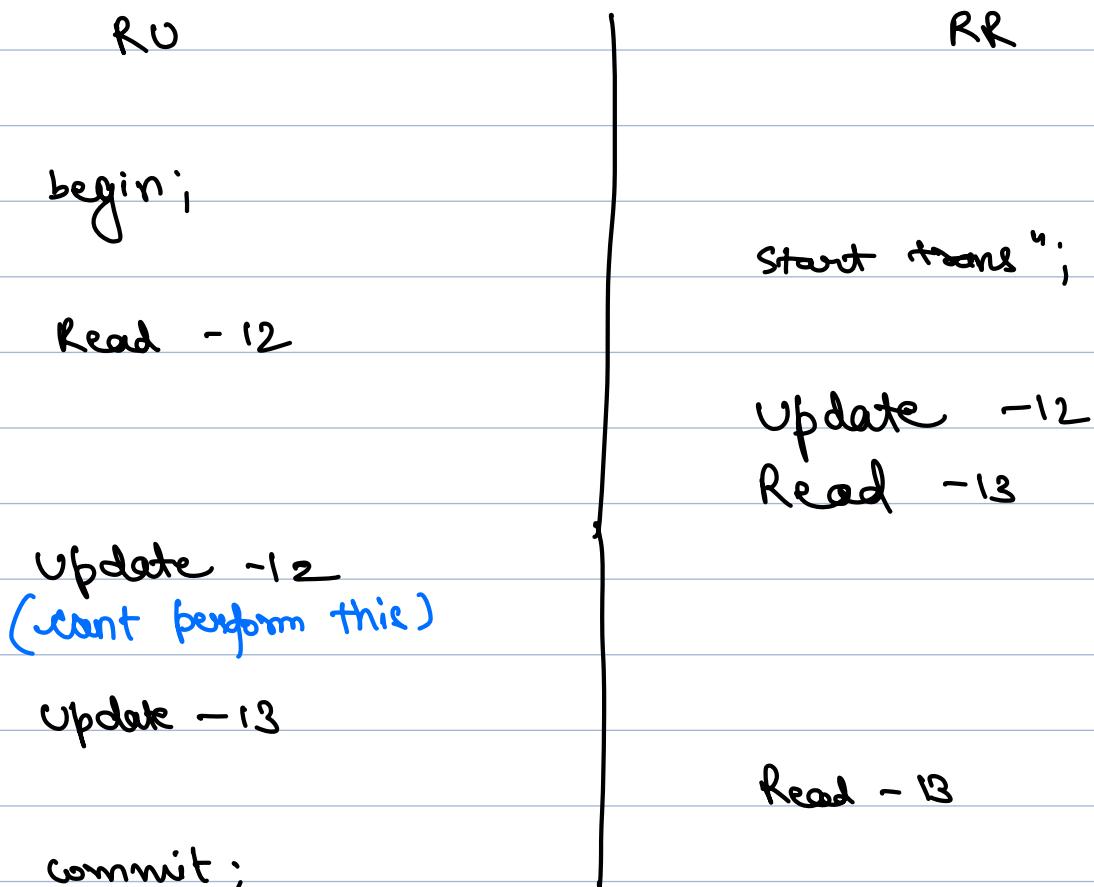
MySQL / SQL DB → by default has auto commit = 1

→ every query executes as a part of new trans.

start transaction ; } begin a "trans" manually → don't need auto commit = 0;

} ==  
==  
==  
commit / roll back }

[CODE].



RR  
Read - 13 (still some value)

commit;

Read - 13 (new value)

## Problem with Repeatable read.

first we learnt Dirty read → RC

then non-repeatable read → RR

solved by

# Repeatable read stores values of which row?

of the row it first read /  
set of rows around it.

(code)

RU

begin

Read - 12

Read - 180

Update - 180

commit

RR

begin

Read - 12

Read - 180 (new value  
whereas  
should have  
been old  
value)

↓

happened bcoz it takes a snapshot of subset (nearby rows)

- Q) Send an email to every user and after sending email set their email-sent = true.

emails		
id	email	email-sent

select email,  
from emails;

# In last psb example  
we assumed, no  
new student is  
added.

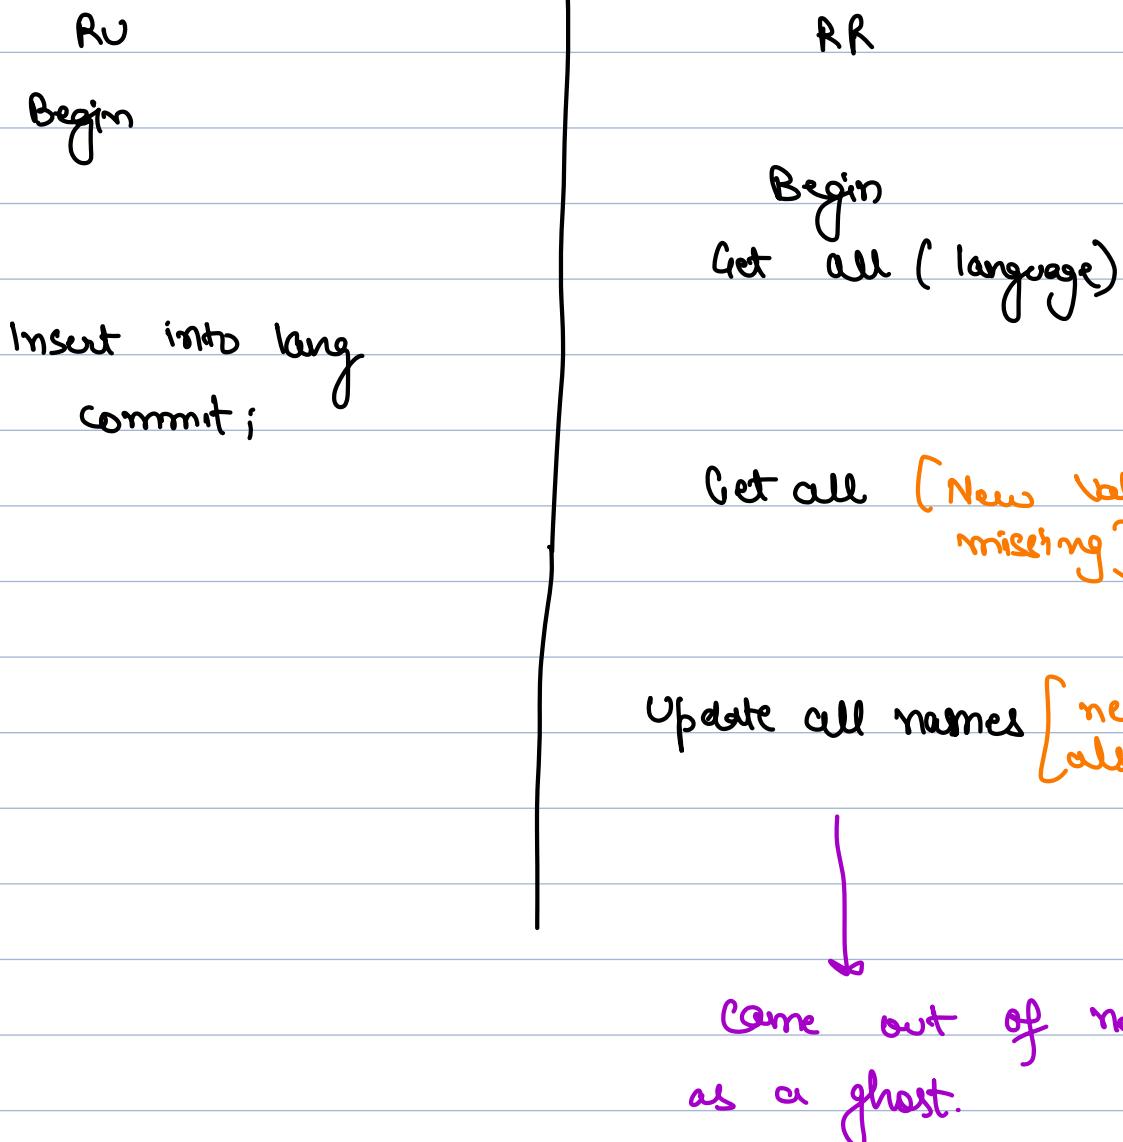
Send emails to them;

One new  
student  
inserted

Update emails  
set email-sent = true;

Repeatable read talks only about rows & their values  
that you initially read. It says nothing about rows  
that you never read.

## [ code ]



{ show above example again with more values }

The problem is known as **Phantom read.**

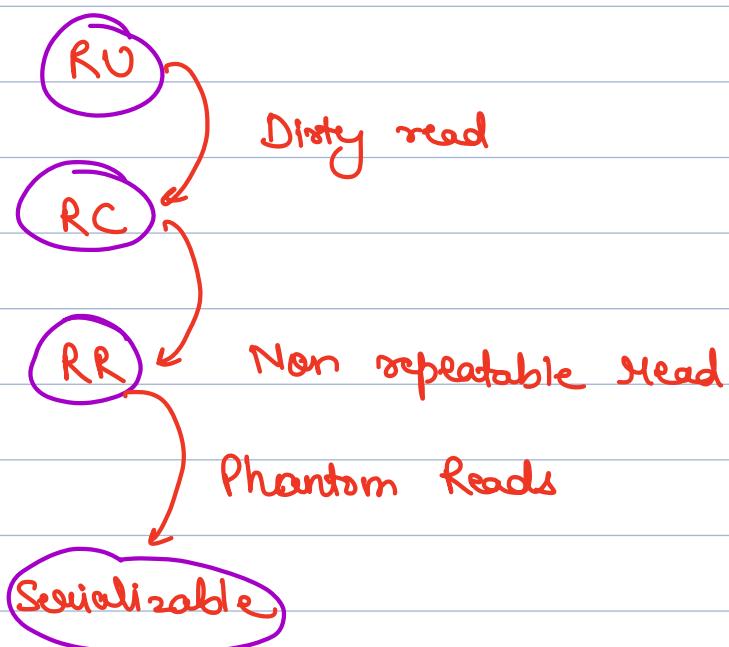
A row comes magically that you had not intended earlier.

# RR only guarantees about the rows that read initially.

→ Why this happens:

Every isolation level only talks about what it will read, not about what/how it will write/update.

\* Within same trans", it will always read your latest value.



# R only guaranteed that rows you had read will remain the same.

# It never guaranteed about the complete table.

- # ① you read rows of the table → Someone inserted.
- ② did some work on rows → that value got updated.
- ③ updated rows of the table

Another example

[CODE]

RU  
begin  
delete & commit

RR  
begin

Read → Meads 10 but  
there are only  
9 in DB.

Update delete value [Nothing happens]  
commit ;

New data will be different.



This is because  
data is not there  
in actual dB.

Serializable.



# If I am doing anything on the table, no one  
else can work on it!

# If somehow → no concurrency, at one time only  
one person can work on one table.

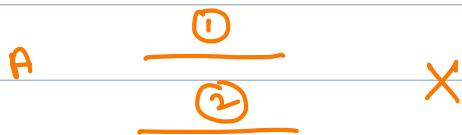
↓  
low performance

Remove concurrency

↳ Banks, BMS.

Serialized → sequential (one after other)

# If a "trans" has started for a row A, the other "trans" can only start after the first trans is completed.



No concurrency | No parallelism

[CODE]

serializable

begin

update

RR

begin  
Read

Read (on updated)

Same  
as  
before

Update (on same item)  
(not allowed)

Serializable

Serializable

Begin  
Read (12 & 13)

select \* from film  
where id in (12, 13)

for update;



I am going to  
update these values,  
don't let anyone  
read them

# If one of them is not serializable, aren't they  
doing harm to themselves by reading values  
they are not sure about.



Thus both need to be serializable.

# Serializable only allows us to read rows that  
are OK to be read.

# In others even if you have a lock, others  
could atleast read the data, Serializable is very  
strict and don't even allow that.

Serializable  $\Rightarrow$  Repeatable Read + Check for lock even for read.



If anyone else has a lock, it doesn't even reach it.

### Example

Book My Show

RR or something  
↑

① Show tickets

② Book tickets

↓  
Serializable

—	—	—	
—			

[Show on website]



One goes one blocked

\* In others you can read any row, in Serializable you can only read rows with No Lock.  
In others which version on data you read is the problem, read/not read was not the problem.

# In serializable  
↓

You can't read → if other is writing or reading for update.

[CODE]

How banks might be working :-

Serializable

Read → film → 12, 13  
for update.

serializable

Read → 12, 13 (Not allow)

# Simply read doesn't take a lock.

# If one person has taken lock for update only then others can't read, else they can.

# Serializable doesn't allow you to even read a row that has a lock.

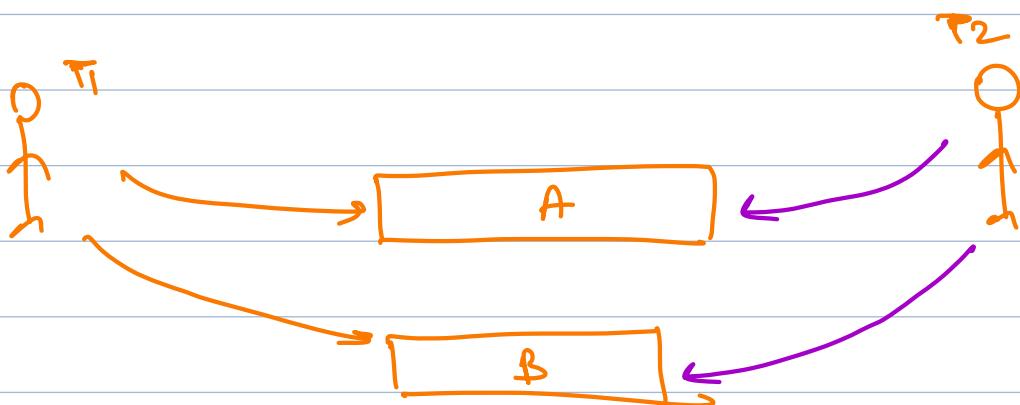
# speed affected but works for bank cases.

# For update → a way to basically say 'I will be writing' soon to this data.

### Deadlocks

# eg → You have crush on her, she has crush on you.  
You are waiting for her to confess & she is waiting for you to confess.

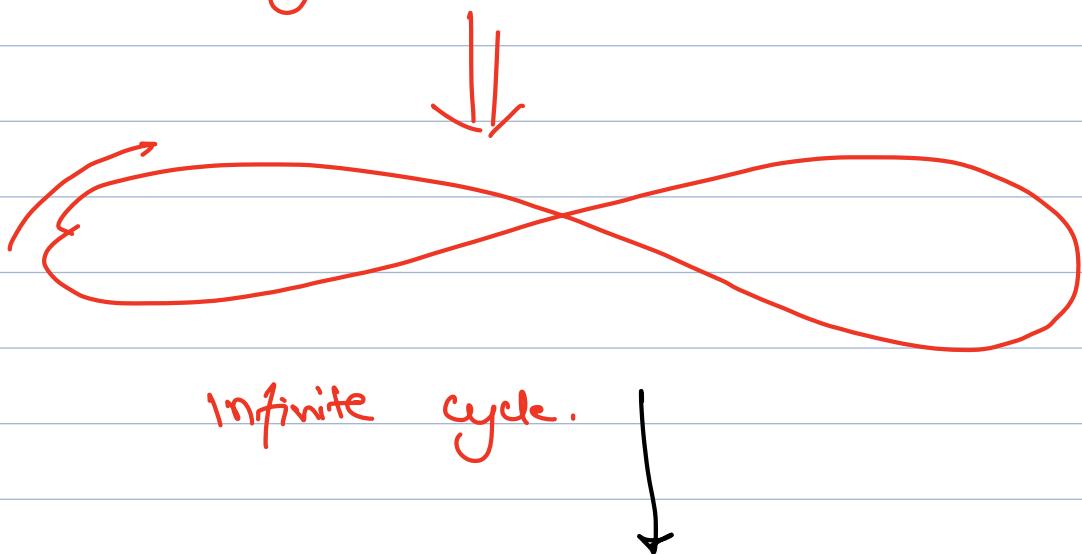
# We learnt in "race" when we write/read for update we take lock on row. (this belongs to me).



$T_1$  will not  
get a lock till  
 $T_2$  completes

$T_2$  will not get  
a lock till  $T_1$   
completely.

#  $T_1$  is waiting for  $T_2$  to complete but  
 $T_2$  is waiting for  $T_1$  to complete



Deadlock  $\rightarrow$  learn later  
in concurrency.

# In OS we will talk how to handle deadlocks.

$\Rightarrow$  SQL detects deadlocks, it automatically kill and  
rolls back one of the trans.

## CODE

begin

✓ Read 12 - for update

\* Read 13 - for update  
(doesn't get lock)  
(this time won't lead to a  
deadlock, just waiting for  
resource)

begin .

✓ Read 13 - for update

Read 12 - for update  
(Deadlock)  
(Trans" cancelled)

